

APÉNDICE C*

Í PIPELININGÍ : conceptos básicos e intermedios

RISC - V

* Computer Architecture , Hennessy & Patterson, 6ta Edición

The load and store instructions in RISC-V

ld x1,80(x2)	Load doubleword	$\text{Regs}[x1] = \text{Mem}[80 + \text{Regs}[x2]]$
lw x1,60(x2)	Load word	$\text{Regs}[x1] \leftarrow_{64} \text{Mem}[60 + \text{Regs}[x2]]_0^{32} \# \# \text{Mem}[60 + \text{Regs}[x2]]$
lwu x1,60(x2)	Load word unsigned	$\text{Regs}[x1] \leftarrow_{64} 0^{32} \# \# \text{Mem}[60 + \text{Regs}[x2]]$
lb x1,40(x3)	Load byte	$\text{Regs}[x1] \leftarrow_{64} (\text{Mem}[40 + \text{Regs}[x3]]_0)^{56} \# \# \text{Mem}[40 + \text{Regs}[x3]]$
lbu x1,40(x3)	Load byte unsigned	$\text{Regs}[x1] \leftarrow_{64} 0^{56} \# \# \text{Mem}[40 + \text{Regs}[x3]]$
lh x1,40(x3)	Load half word	$\text{Regs}[x1] \leftarrow_{64} (\text{Mem}[40 + \text{Regs}[x3]]_0)^{48} \# \# \text{Mem}[40 + \text{Regs}[x3]]$
flw f0,50(x3)	Load FP single	$\text{Regs}[f0] \leftarrow_{64} \text{Mem}[50 + \text{Regs}[x3]] \# \# 0^{32}$
fld f0,50(x2)	Load FP double	$\text{Regs}[f0] \leftarrow_{64} \text{Mem}[50 + \text{Regs}[x2]]$
sd x2,400(x3)	Store double	$\text{Mem}[400 + \text{Regs}[x3]] \leftarrow_{64} \text{Regs}[x2]$
sw x3,500(x4)	Store word	$\text{Mem}[500 + \text{Regs}[x4]] \leftarrow_{32} \text{Regs}[x3]_{32..63}$
fsw f0,40(x3)	Store FP single	$\text{Mem}[40 + \text{Regs}[x3]] \leftarrow_{32} \text{Regs}[f0]_{0..31}$
fsd f0,40(x3)	Store FP double	$\text{Mem}[40 + \text{Regs}[x3]] \leftarrow 64 \text{Regs}[f0]$
sh x3,502(x2)	Store half	$\text{Mem}[502 + \text{Regs}[x2]] \leftarrow_{16} \text{Regs}[x3]_{48..63}$
sb x2,41(x3)	Store byte	$\text{Mem}[41 + \text{Regs}[x3]] \leftarrow_8 \text{Regs}[x2]_{56..63}$

The basic ALU instructions in RISC-V

add x1,x2,x3	Add	Regs[x1] Regs[x2]+Regs[x3]
addi x1,x2,3	Add immediate unsigned	Regs[x1] Regs[x2]+3
lui x1,42	Load upper immediate	Regs[x1] = 0 ³² ##42##0 ¹²
sll x1,x2,5	Shift left logical	Regs[x1] = Regs[x2]<<5
slt x1,x2,x3	Set less than	if (Regs[x2]<Regs[x3]) Regs[x1] = 1 else Regs[x1] = 0

Typical control flow instructions in RISC-V

jal x1,offset	Jump and link	Regs[x1]= PC+4; PC= PC + (offset<<1)
jalr x1,x2,offset	Jump and link register	Regs[x1]= PC+4; PC = Regs[x2]+offset
beq x3,x4,name	Branch equal	if (Regs[x3]==Regs[x4]) PC = PC + (offset<<1)
bgt x3,x4,name	Branch greater than	if (Regs[x3]>Regs[x4]) PC= PC + (offset<<1)

```
if (aa==2)
    aa=0;
if (bb==2)
    bb=0;
if (aa!=bb) {
```

Here is the RISC-V code that we would typically generate for this code fragment assuming that `aa` and `bb` are assigned to registers `x1` and `x2`:

```
addi  x3,x1,-2
bnez x3,L1      //branch b1  (aa!=2)
add   x1,x0,x0  //aa=0
L1: addi  x3,x2,-2
bnez x3,L2      //branch b2  (bb!=2)
add   x2,x0,x0  //bb=0
L2: sub   x3,x1,x2 //x3=aa-bb
beqz x3,L3      //branch b3  (aa==bb)
```

31	25 24	20 19	15 14 12 11	7 6	0
funct7	rs2	rs1	funct3	rd	opcode
imm [11:0]	rs1	funct3	rd	opcode	I-type
imm [11:5]	rs2	rs1	funct3	imm [4:0]	opcode
imm [12]	imm [10:5]	rs2	rs1	funct3	imm [4:1 11] opcode
imm [31:12]				rd	opcode
imm [20 10:1 11 19:12]				rd	opcode

Figura 1.7 Formatos de codificación base para RISC-V Todas son de 32 bits .

Tipo **R**: Para operaciones de enteros R-R

Tipo **I**: Para Loads y operaciones con inmediato.

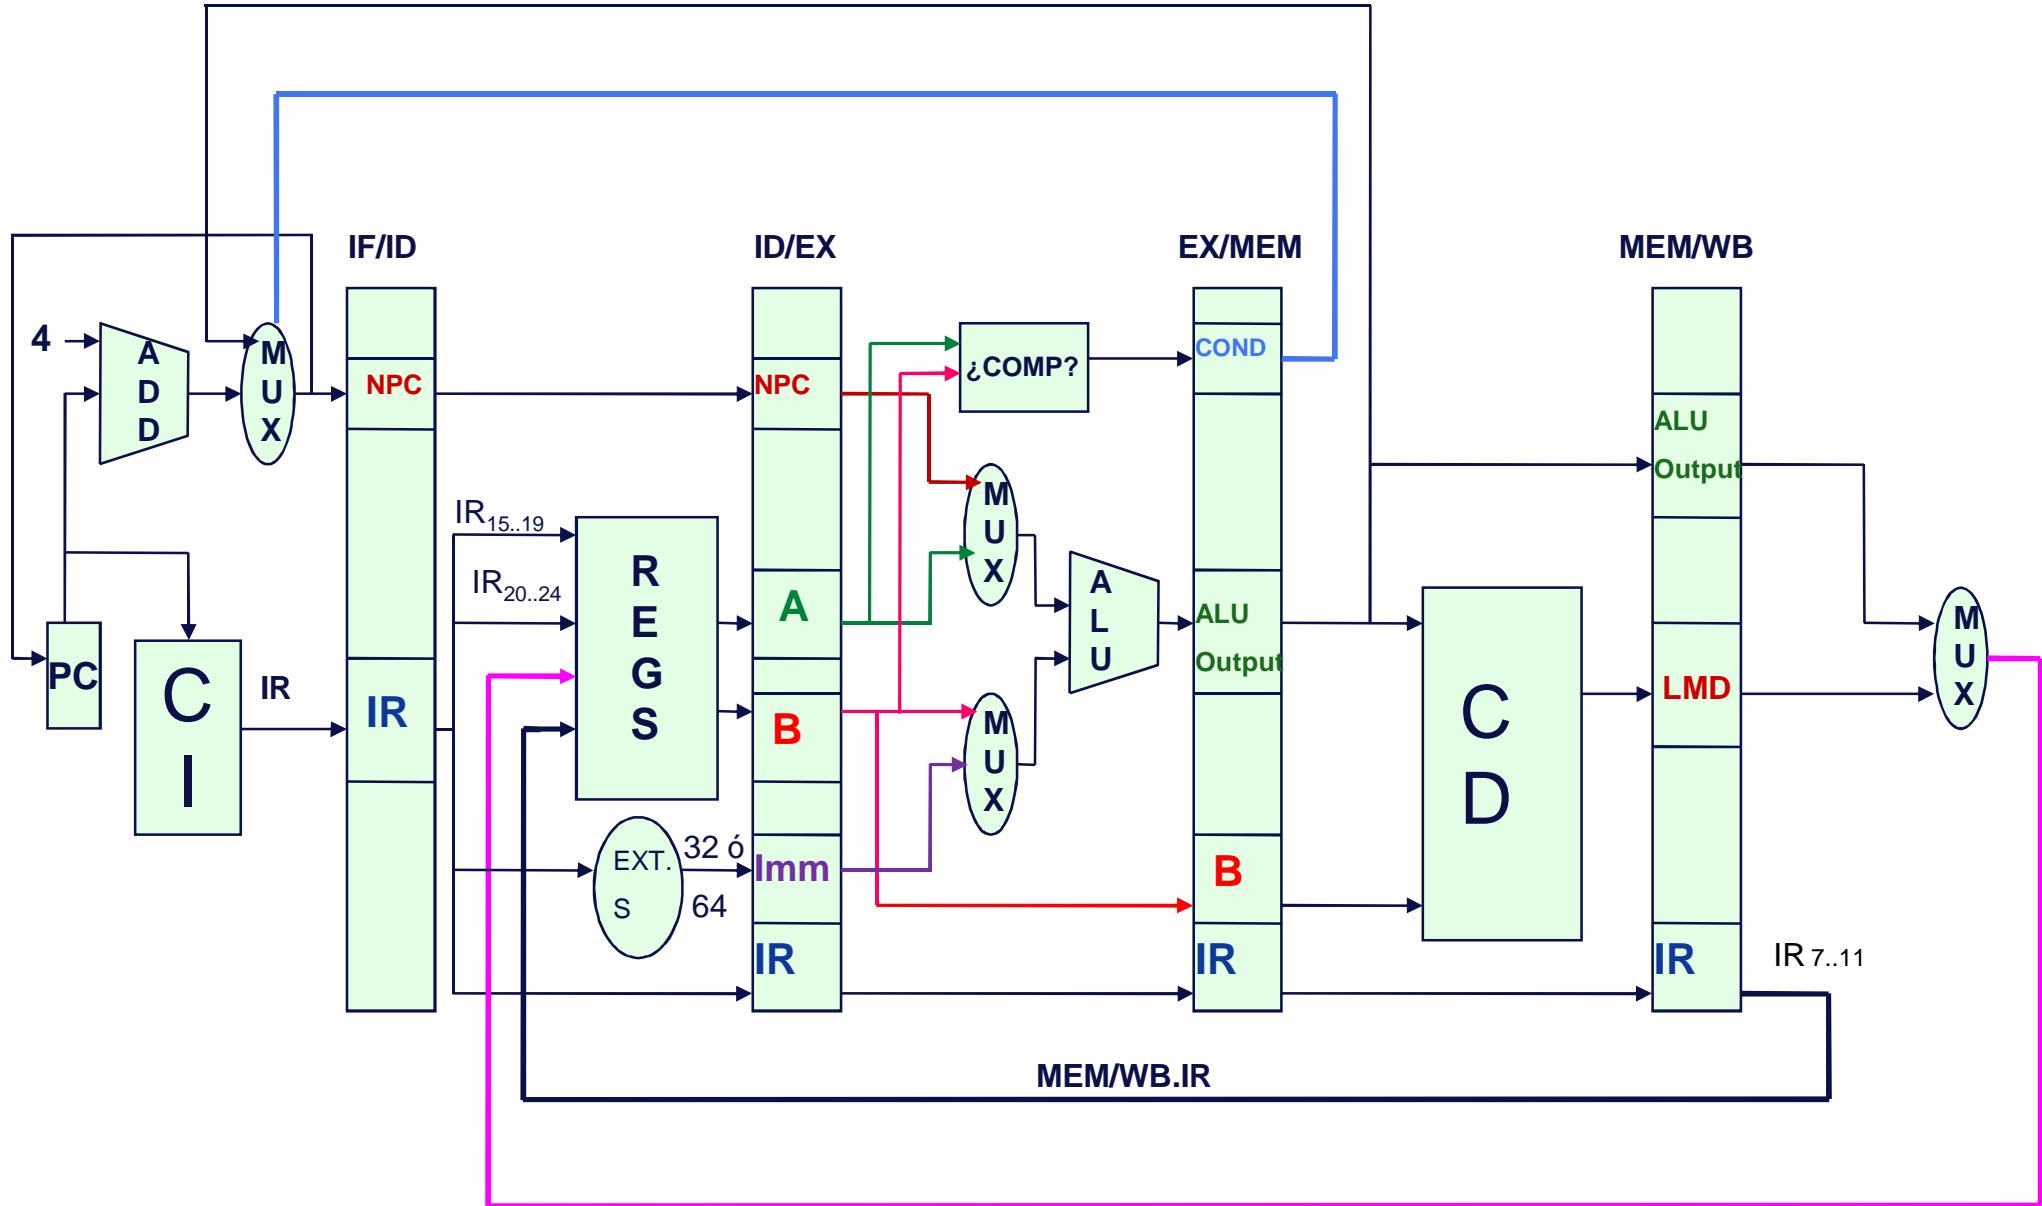
Tipo **B**: Para Branches

Tipo **J**: Para Jumps and link

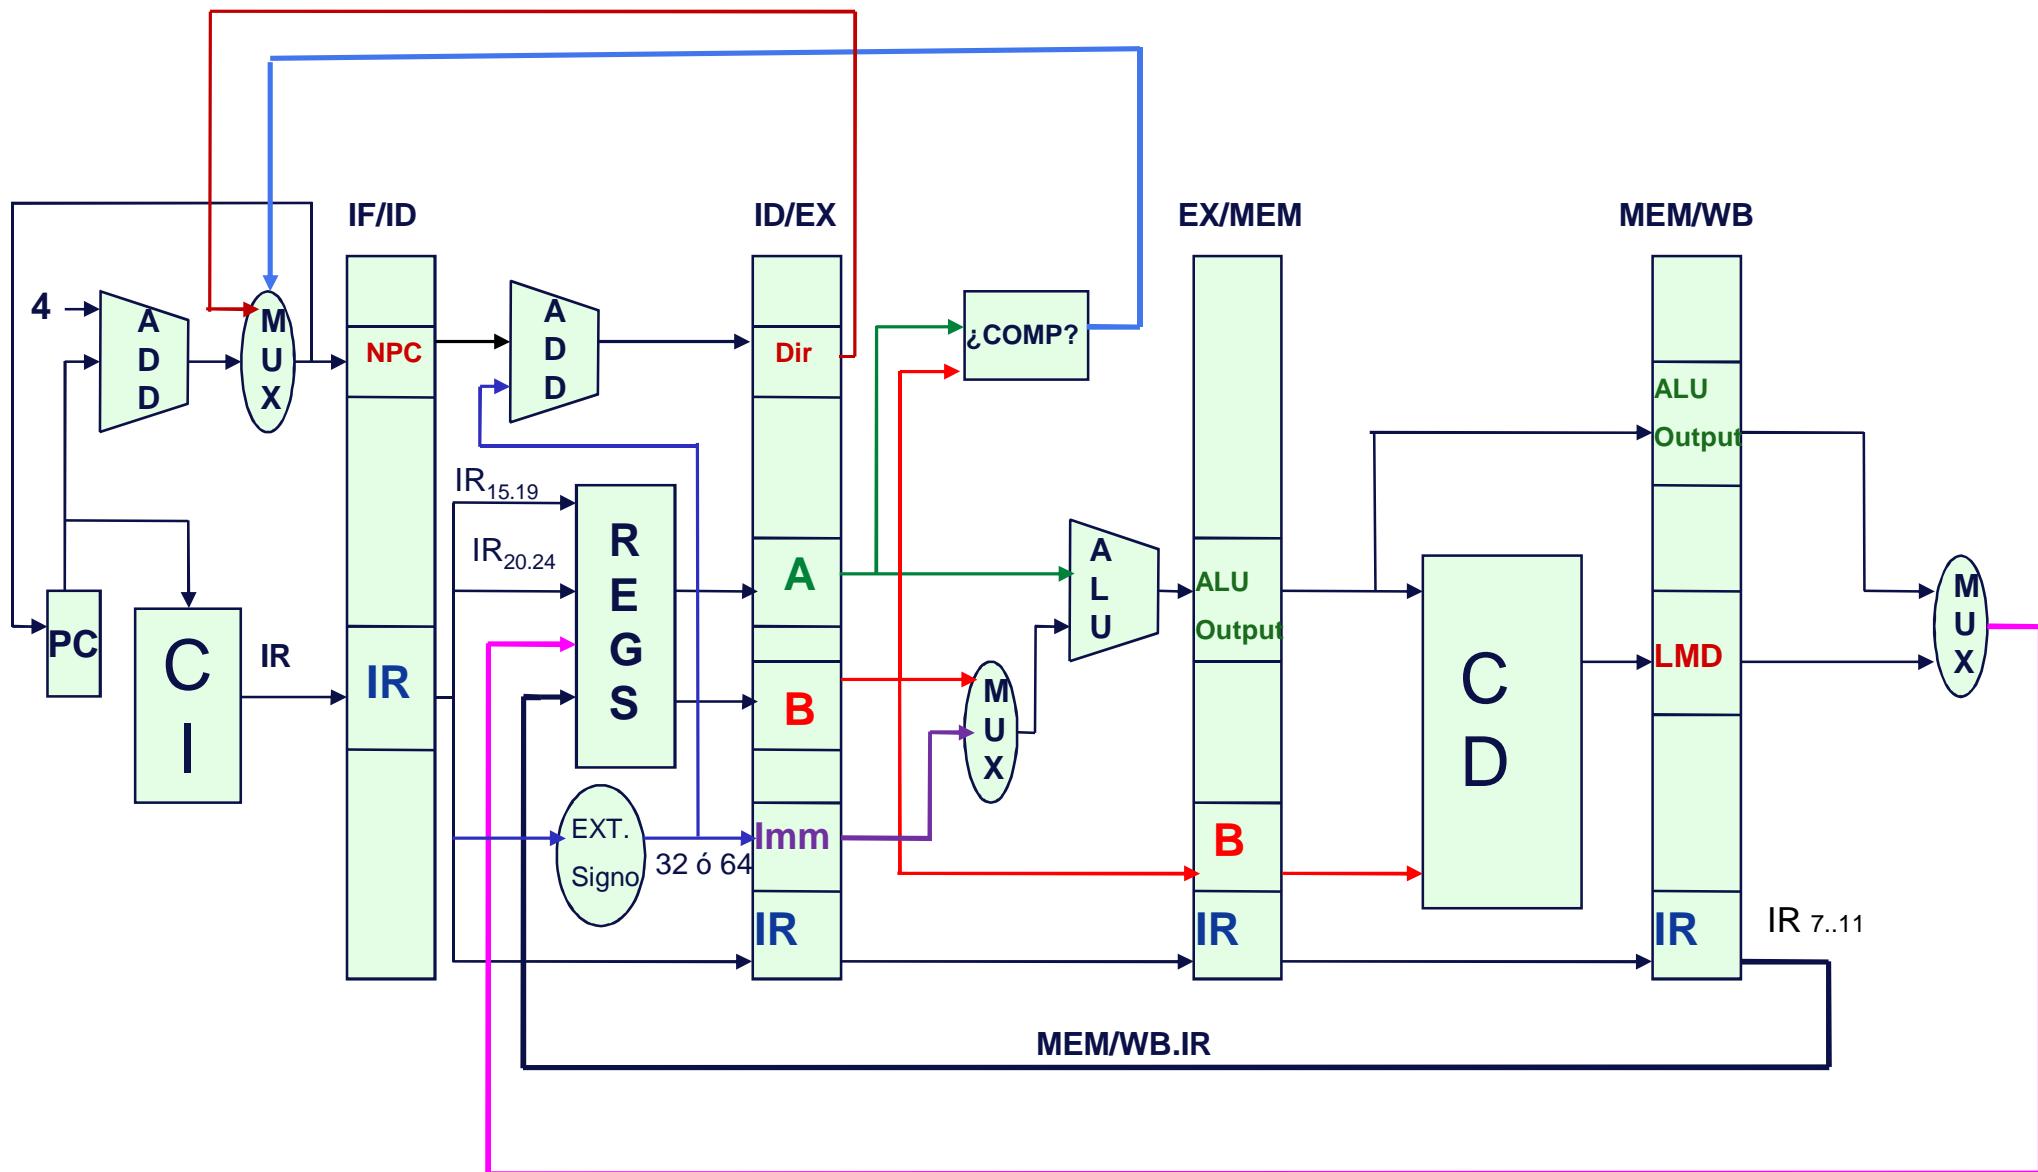
Tipo **S**: Para Stores (eso para que los especificadores de los 3 registros estén siempre en la misma posición - (rd, rs1, rs2)

Tipo **U**: Para instrucciones con %inmediato ancho+(LUI, AUIPC)

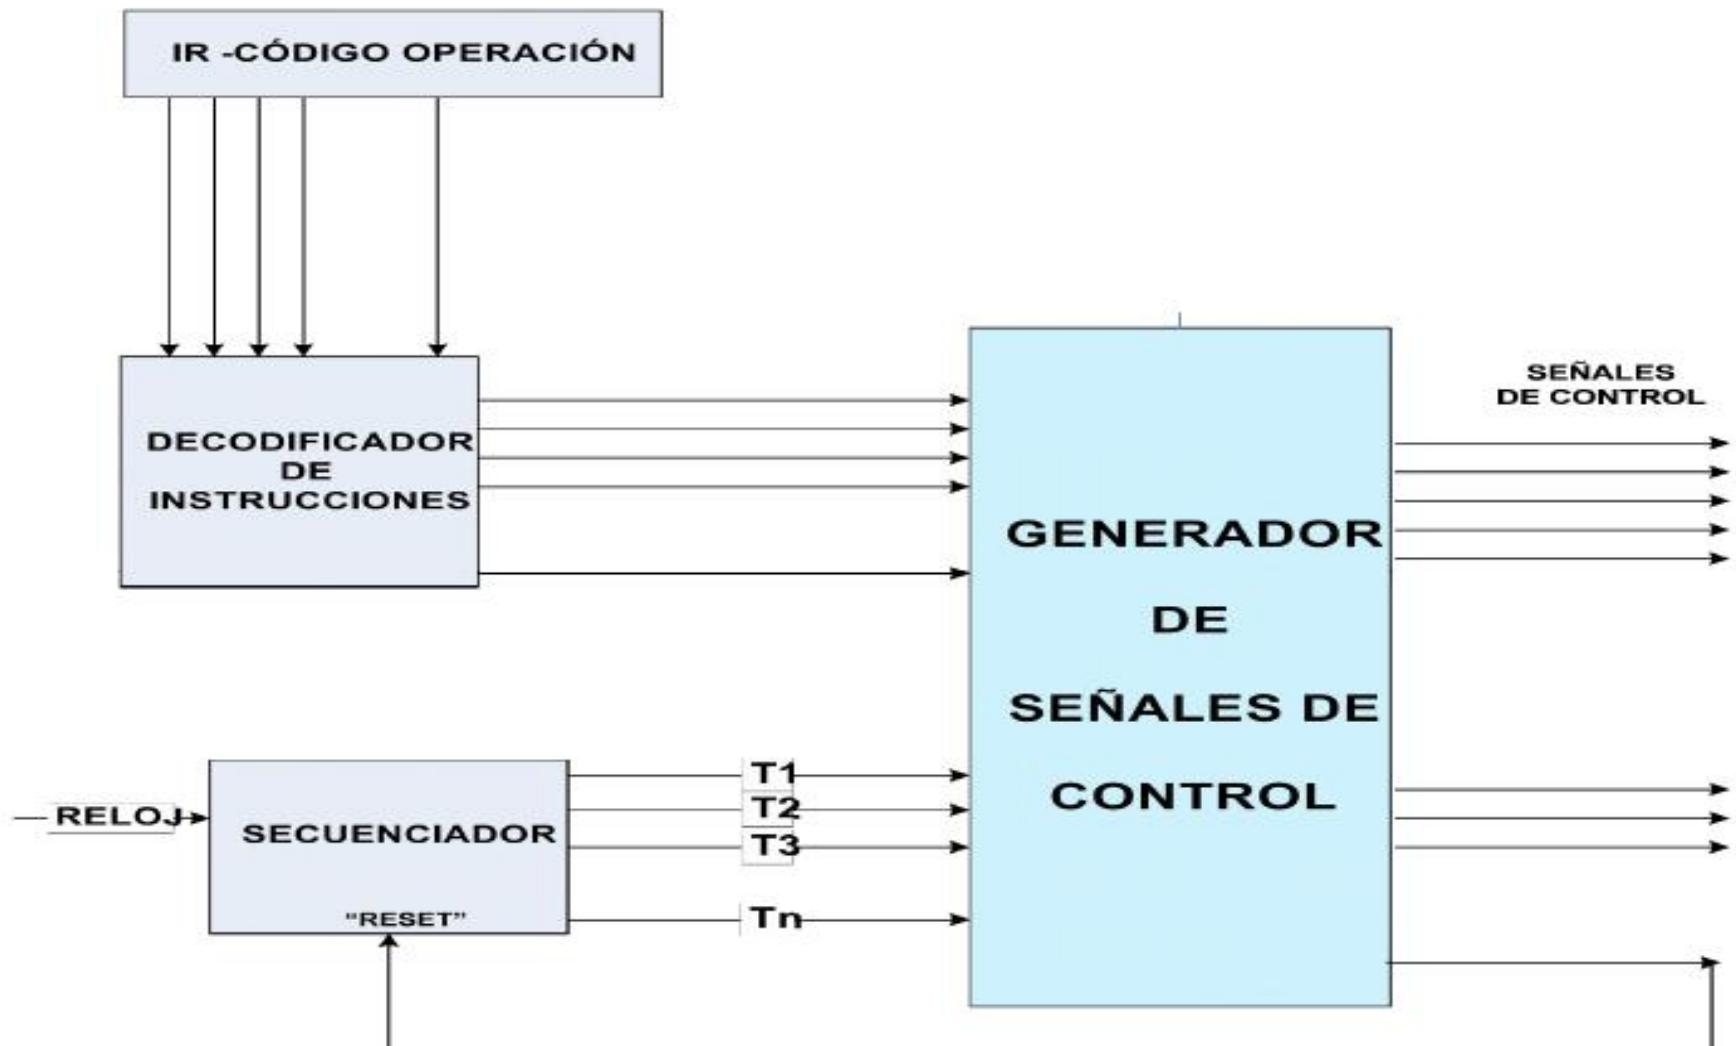
PIPELINE ENTEROS RISC-V



PIPELINE ENTEROS RISC-V MODIFICADO



CONTROL ALAMBRADO (Hardwired)



EJECUCIÓN EN PIPELINE

TABLA DE TIEMPOS

Explicación	Instrucción	1	2	3	4	5	6	7	8	9
$x1 = M[60 + x2]$	<u>lw x1,60(x2)</u>	IF	ID	EX	M	WB				
$x7 = x4 + x5$	<u>add x7,x4,x5</u>		IF	ID	EX	M	WB			
$x8 = x9 - x6$	<u>sub x8,x9,x6</u>			IF	ID	EX	M	WB		
$x3 = x10 + 6$	<u>addi x3,x10,6</u>				IF	ID	EX	M	WB	
$M[500 + x4] = x12$	<u>sw x12,500(x4)</u>					IF	ID	EX	M	WB

CONFLICTOS DE DATOS

SOLUCIÓN 1- DETENER PIPELINE

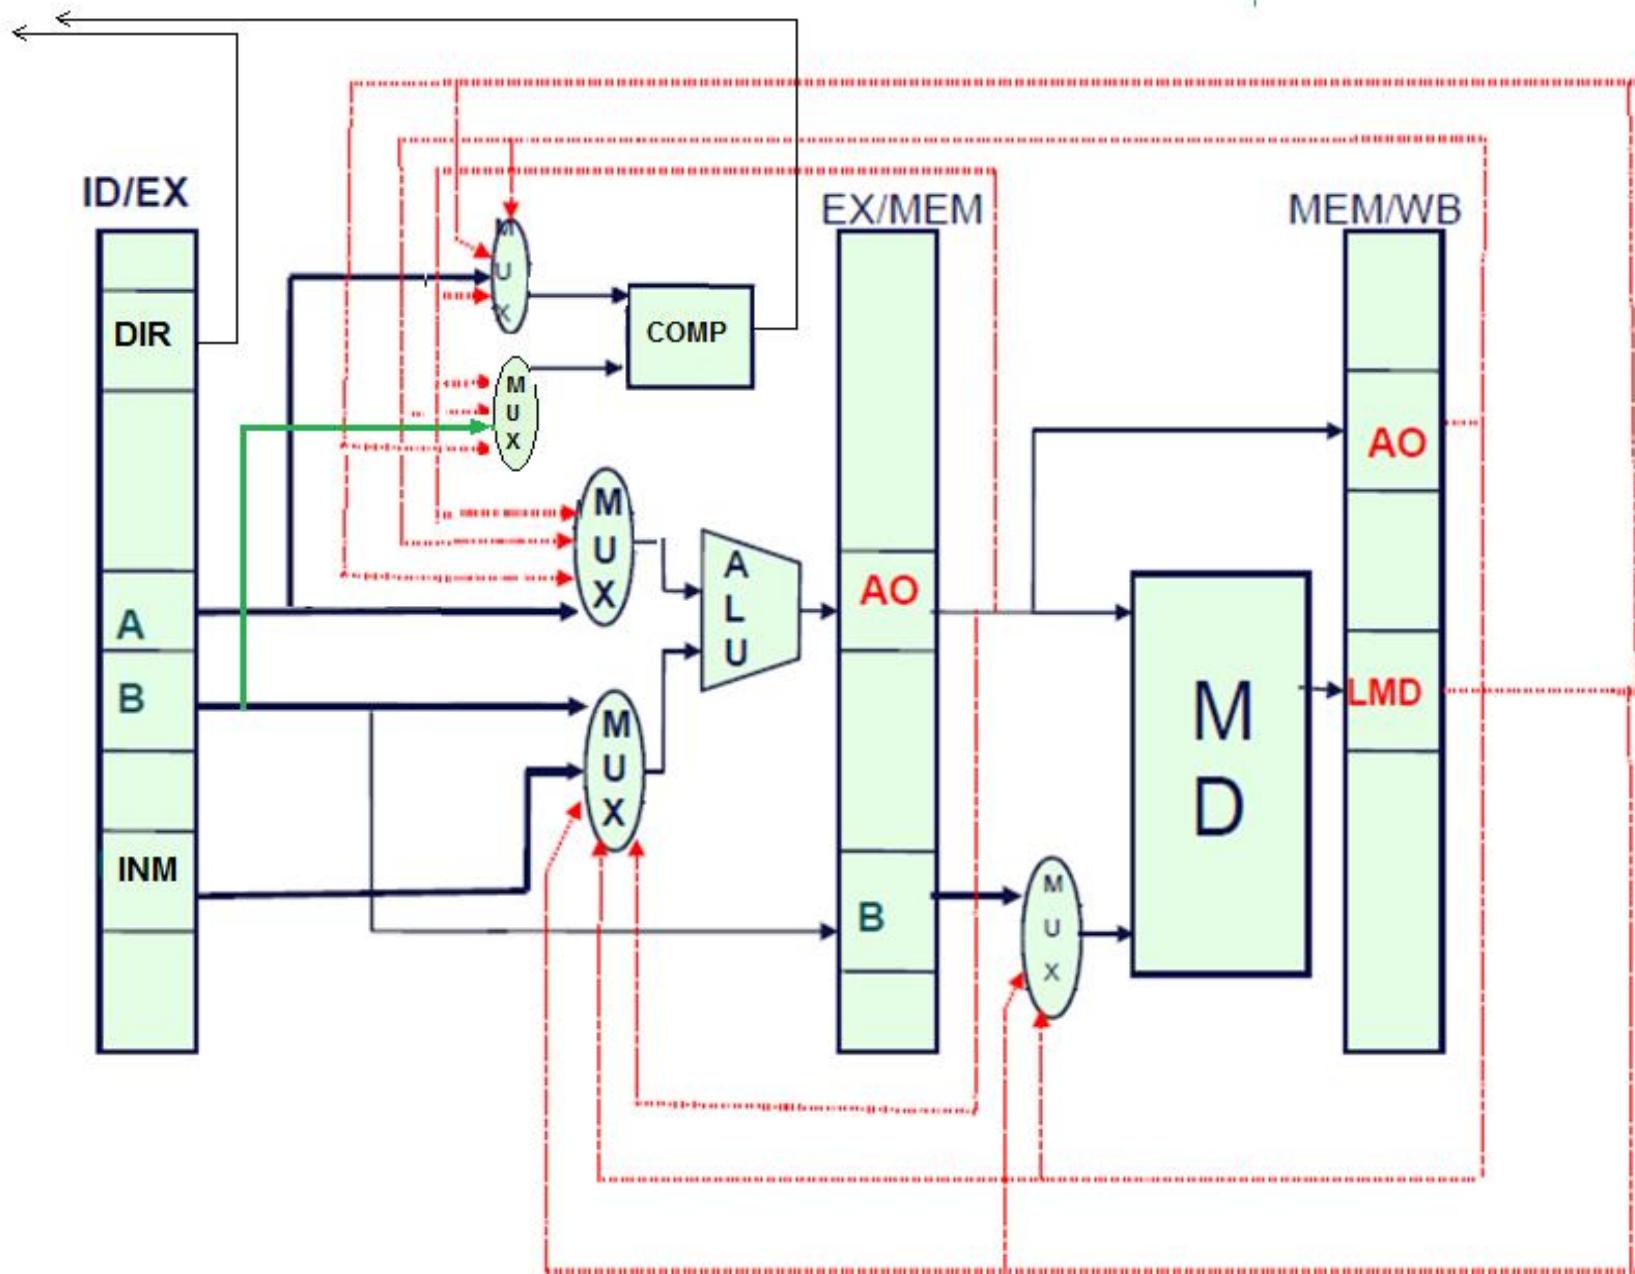
	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
add x2,x4,x5	IF	ID	EX	M	WB										
lw x1,60(x2)	IF	ID	ID	ID	EX	M	WB								
addi x3,x1,6				IF	ID	ID	ID	EX	M	WB					
sw x3,500(x1)						IF	ID	ID	ID	EX	M	WB			
sub x8,x1,x2									IF	ID	EX	M	WB		

CONFLICTOS DE DATOS

SOLUCIÓN 2- FORWARDING

	1	2	3	4	5	6	7	8	9	10	11
add x2,x4,x5	IF	ID	EX	M	WB						
lw x1,60(x2)		IF	ID	EX	M	WB					
addi x3,x1,6			IF	ID	stall	EX	M	WB			
sw x3,500(x1)				IF	stall	ID	EX	M	WB		
sub x8,x1,x2					stall	IF	ID	EX	M	WB	

Forwarding requiere 3 entradas extra en multiplexores del ALU, de COMP.
y a mux de MD



Hacer aquí:

Ejemplos y práctica con FORWARDING

CONFLICTOS DE CONTROL

add x1,x2,x3

beq x12,x0,skip

sub x4,x5,x6

add x5,x4,x9

skip: or x7,x8,x9

ETC...

SOLUCIÓN 1: DETENER PIPELINE

SIN PREDICCIÓN - FALSO

	1	2	3	4	5	6	7	8	9	10	11	CPI
add x1,x2,x3	IF	ID	EX	M	WB							1
beq x12,x0,skip		IF	ID	EX	M	WB						1
sub x4,x5,x6			IF									
sub x4,x5,x6					IF	ID	EX	M	WB			3
add x5,x4,x9						IF	ID	EX	M	WB		1
skip: or x7,x8,x9							IF	ID	EX	M	WB	1
ETC...								IF	ID	EX	M	WB

SIN PREDICCIÓN - VERDADERO

SOLUCIÓN 2 (para disminuir ciclos de retraso): PREDICCIÓN NO TOMADO (%ESTÁTICA+)

PREDICCIÓN NO TOMADO - FALSO

PREDICCIÓN NO TOMADO - VERDADERO

SOLUCIÓN 3 (dism. retraso): BRANCH RETRASADO (Delayed Branch)

Compilador incluye instrucciones después del branch

Mientras se resuelve el branch, el procesador las comienza a ejecuta, e independientemente del resultado de este, concluye su ejecución

Estas instrucciones:

- se toman del código anterior al branch,
- no dependen del resultado del branch,
- no afectan su condición y
- no tienen conflictos entre ellas

BRANCH RETRASADO - EJEMPLO

X y V dos vectores enteros de longitud 32

X se encuentra en memoria a partir del byte 0 y
V a partir del byte 128)

```
for (int i = 0; i < 32, i++) {
```

```
    V[31-i] = X[i] + V[31-i];
```

```
}
```

Se resuelve sin branch retrasado:

x1 tiene dirección del primer elemento de X, es decir, contiene un 0

x2 del último de V, o sea, tiene el valor 252

x3 es 32, la dimensión de X y de V:

AQUÍ: **lw** x8, 0(x2)
 lw x9, 0(x1)
 add x8, x8, x9
 sw x8, 0(x2)
 addi x2, x2, - 4
 addi x1, x1, 4
 addi x3, x3, -1
 bnez x3, AQUÍ

Se resuelve con branch retrasado se incluyen 2 instrucciones buenas:

Código anterior (original)

AQUÍ: lw x8, 0(x2)
lw x9, 0(x1)
add x8, x8, x9
sw x8, 0(x2)
addi x2, x2, -4
addi x1, x1, 4
addi x3, x3, -1
bnez x3, AQUÍ

Código nuevo por branch retrasado

AQUÍ: lw x8, 0(x2)
lw x9, 0(x1)
add x8, x8, x9
addi x1, x1, 4
addi x3, x3, -1
bnez x3, AQUÍ
sw x8, 0(x2)
addi x2, x2, -4

SOLUCIÓN 4: PREDICCIÓN DINAMICA

(Predictores y Branch-Prediction Buffers)

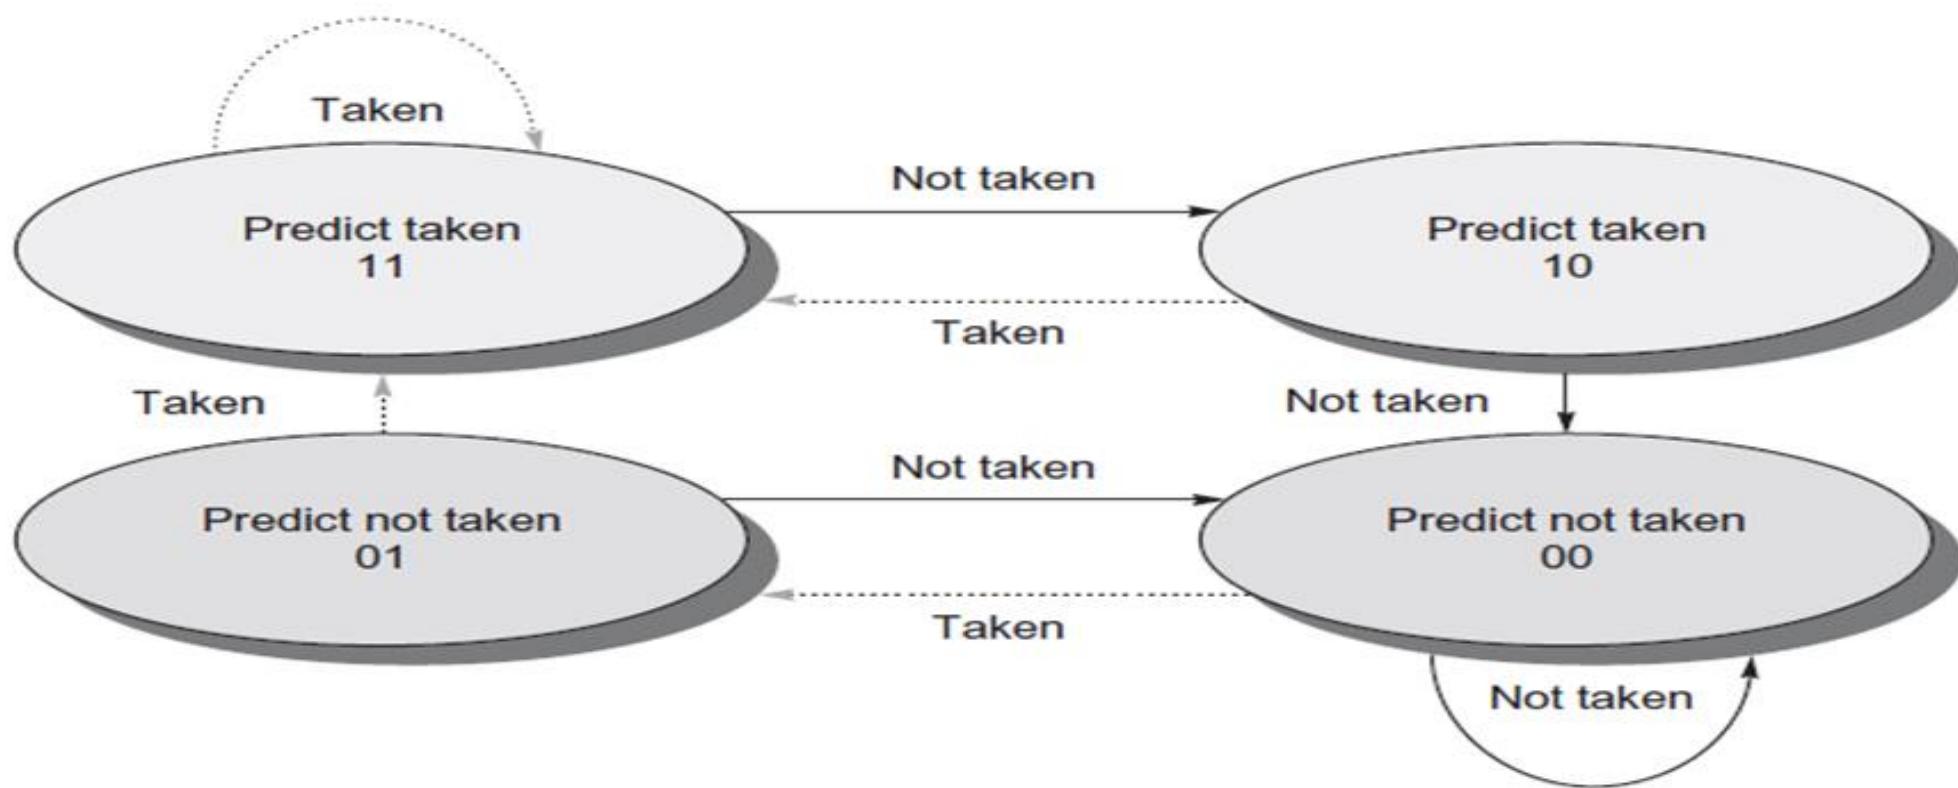
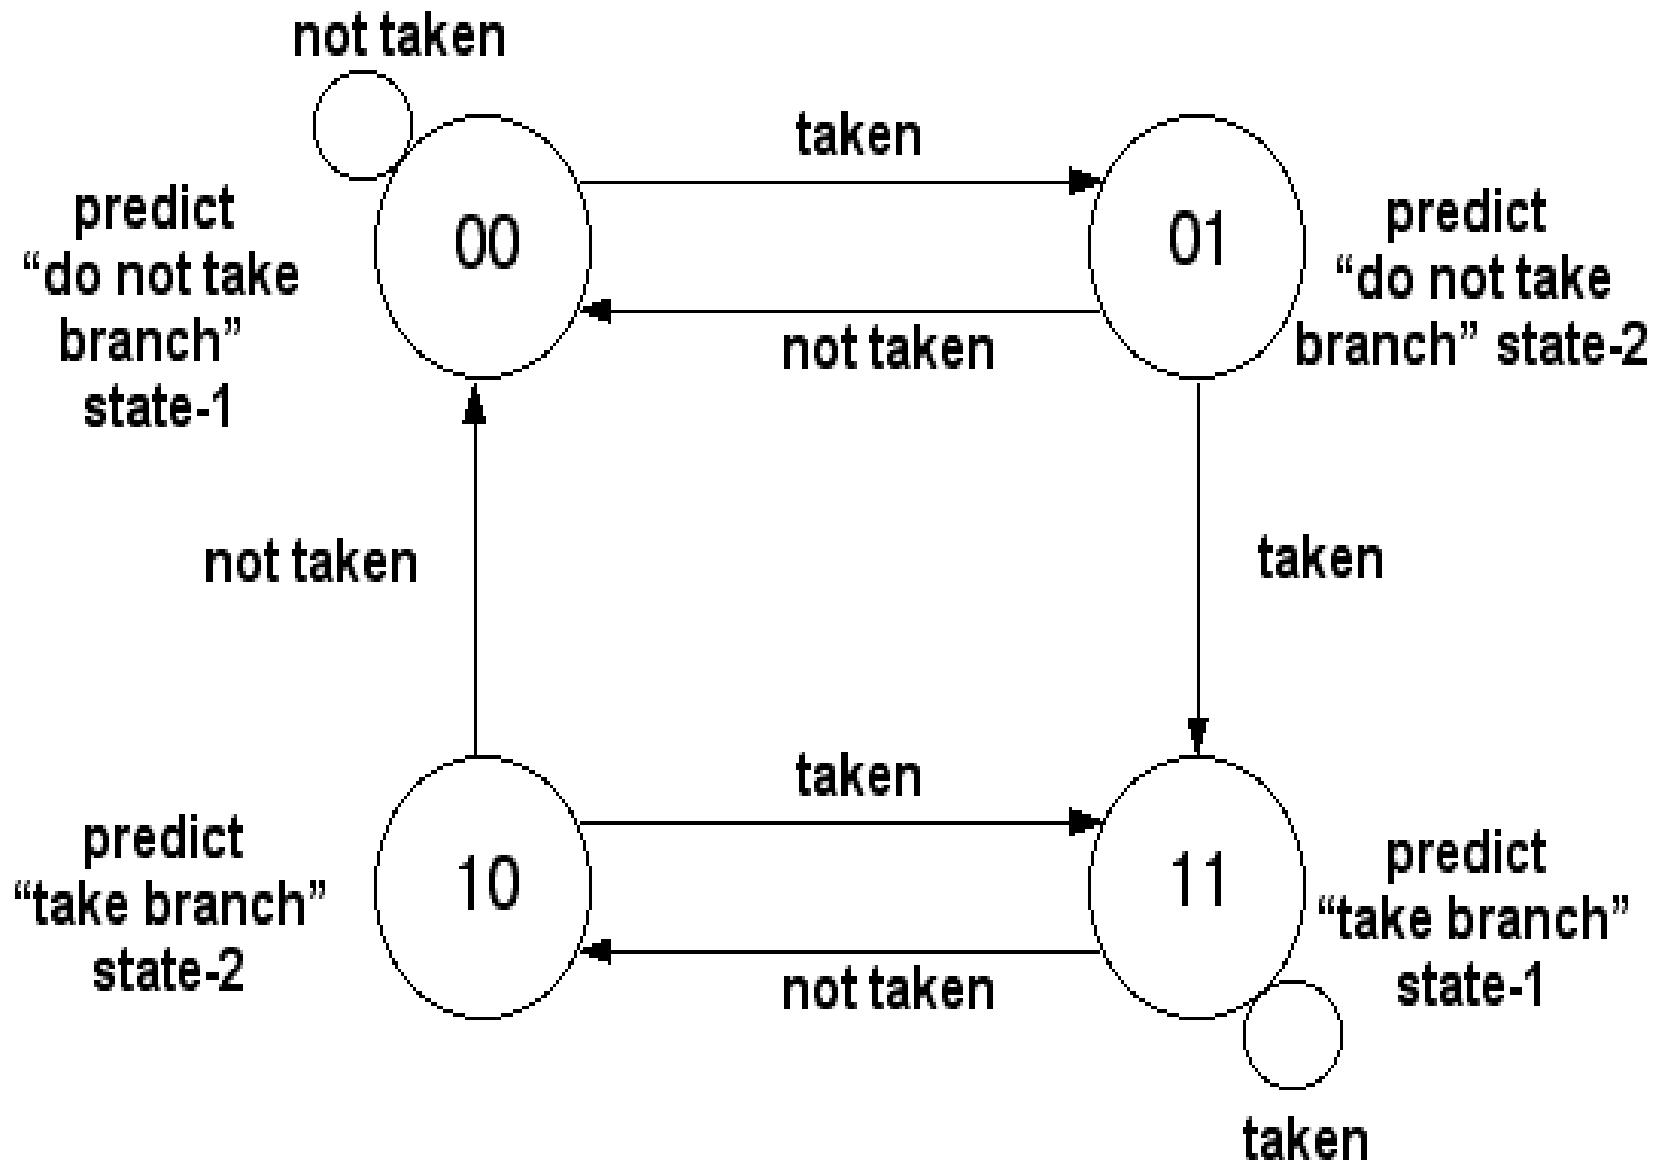
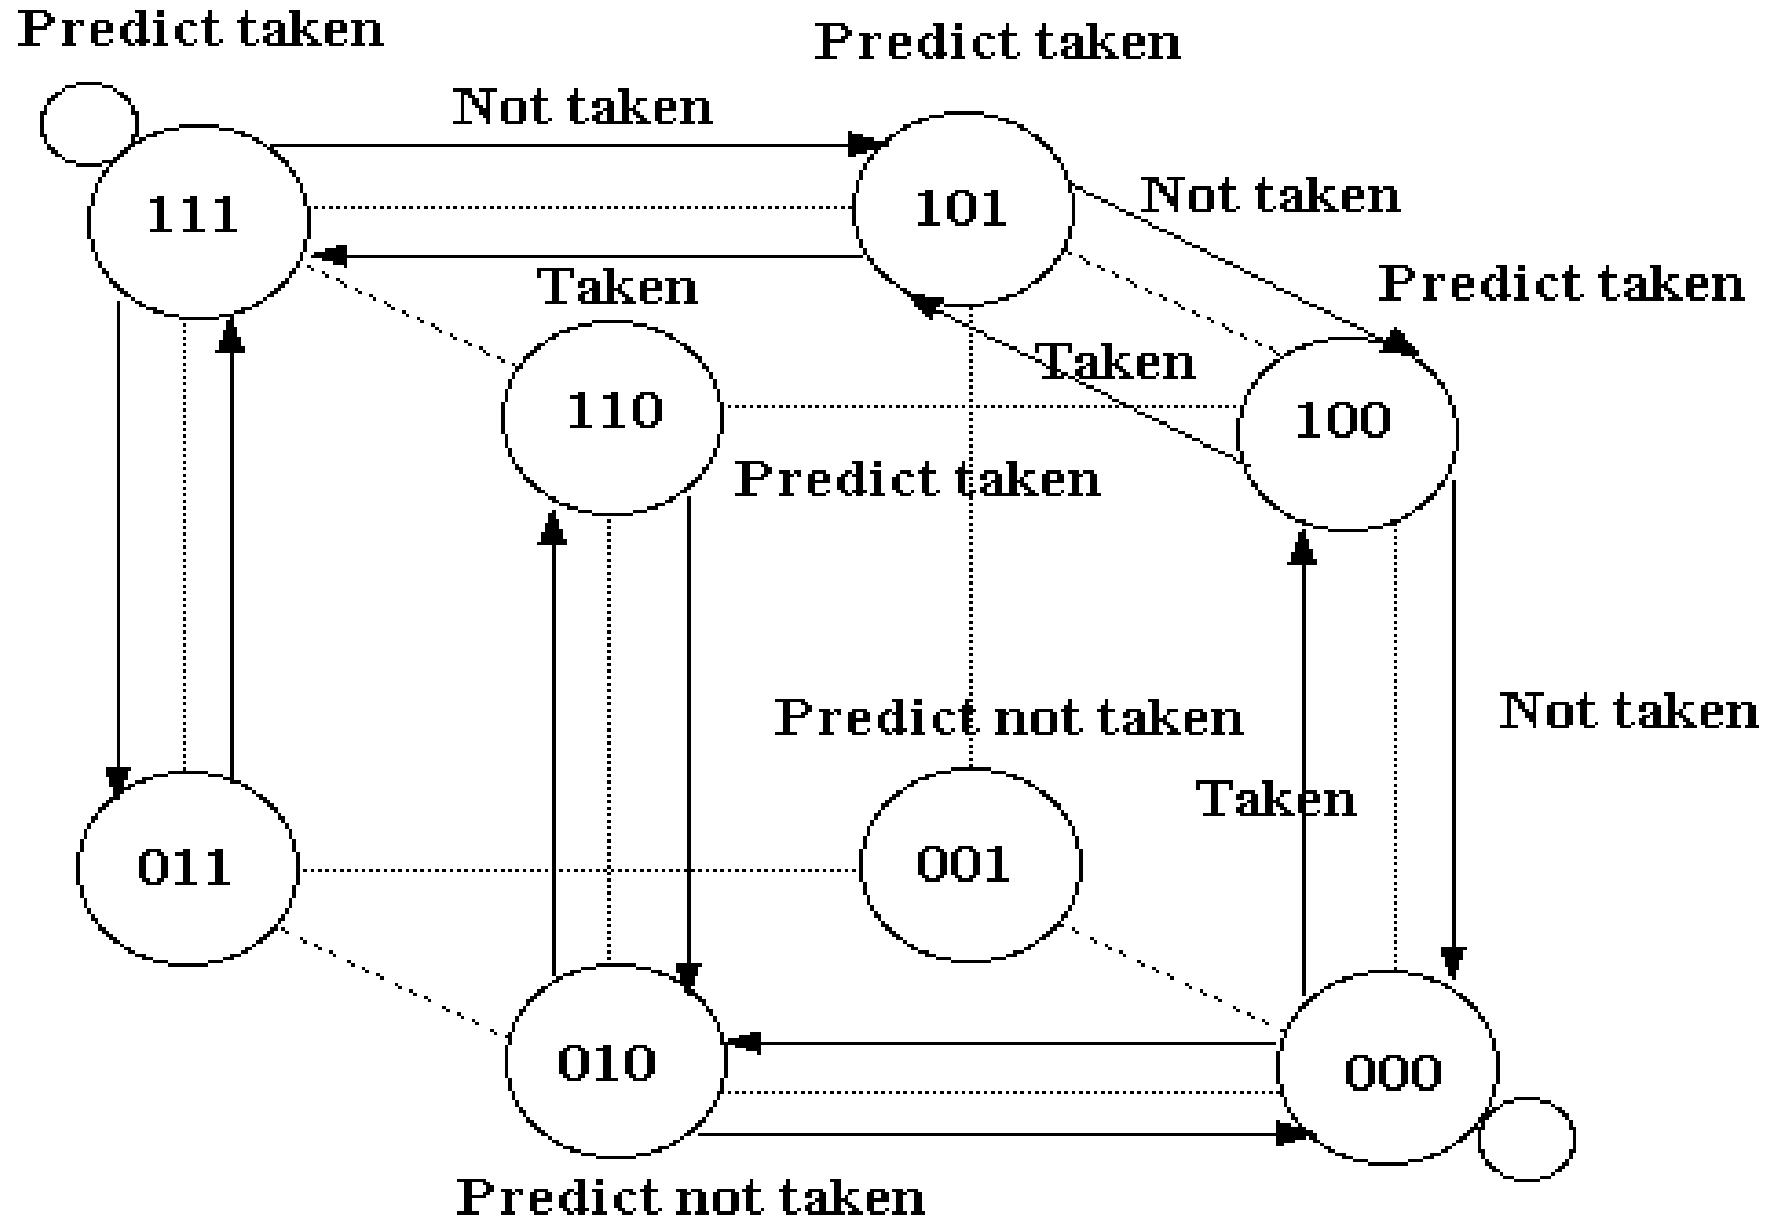


Figure C.15 The states in a 2-bit prediction scheme. By using 2 bits rather than 1, a branch that strongly favors taken or not taken—as many branches do—will be mispredicted less often than with a 1-bit predictor. The 2 bits are used to encode the four states in the system. The 2-bit scheme is actually a specialization of a more general scheme that has an n -bit saturating counter for each entry in the prediction buffer. With an n -bit counter, the counter can take on values between 0 and $2^n - 1$: when the counter is greater than or equal to one-half of its maximum value ($2^n - 1$), the branch is predicted as taken; otherwise, it is predicted as untaken. Studies of n -bit predictors have shown that the 2-bit predictors do almost as well, thus most systems rely on 2-bit branch predictors rather than the more general n -bit predictors.

PREDICTOR DE 2 BITS



PREDICTOR DE 3 BITS



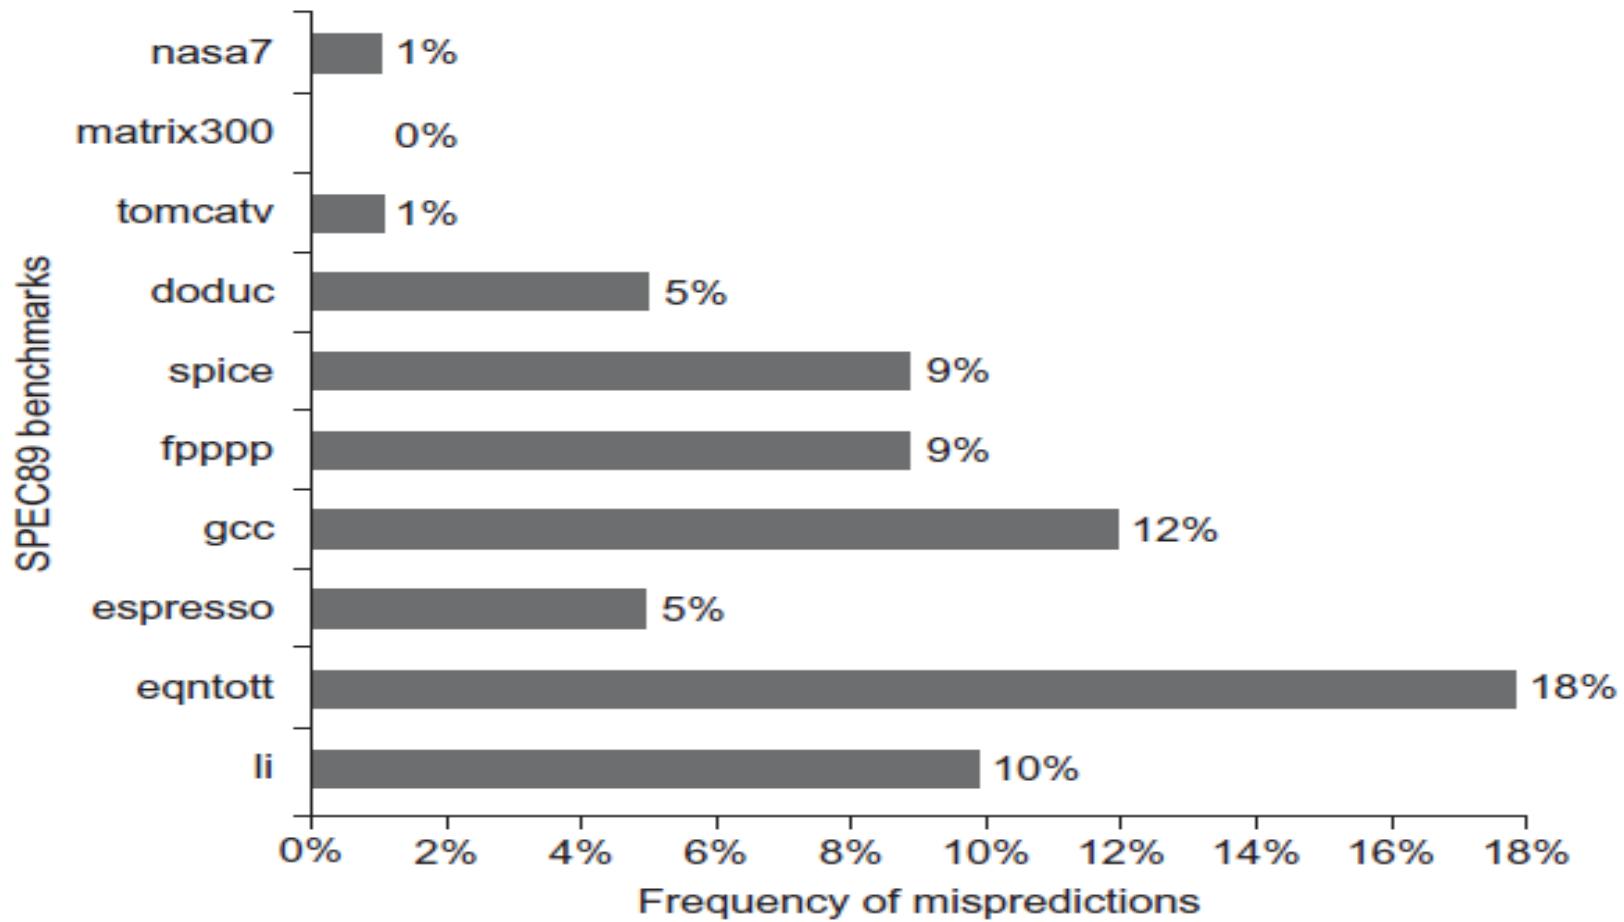
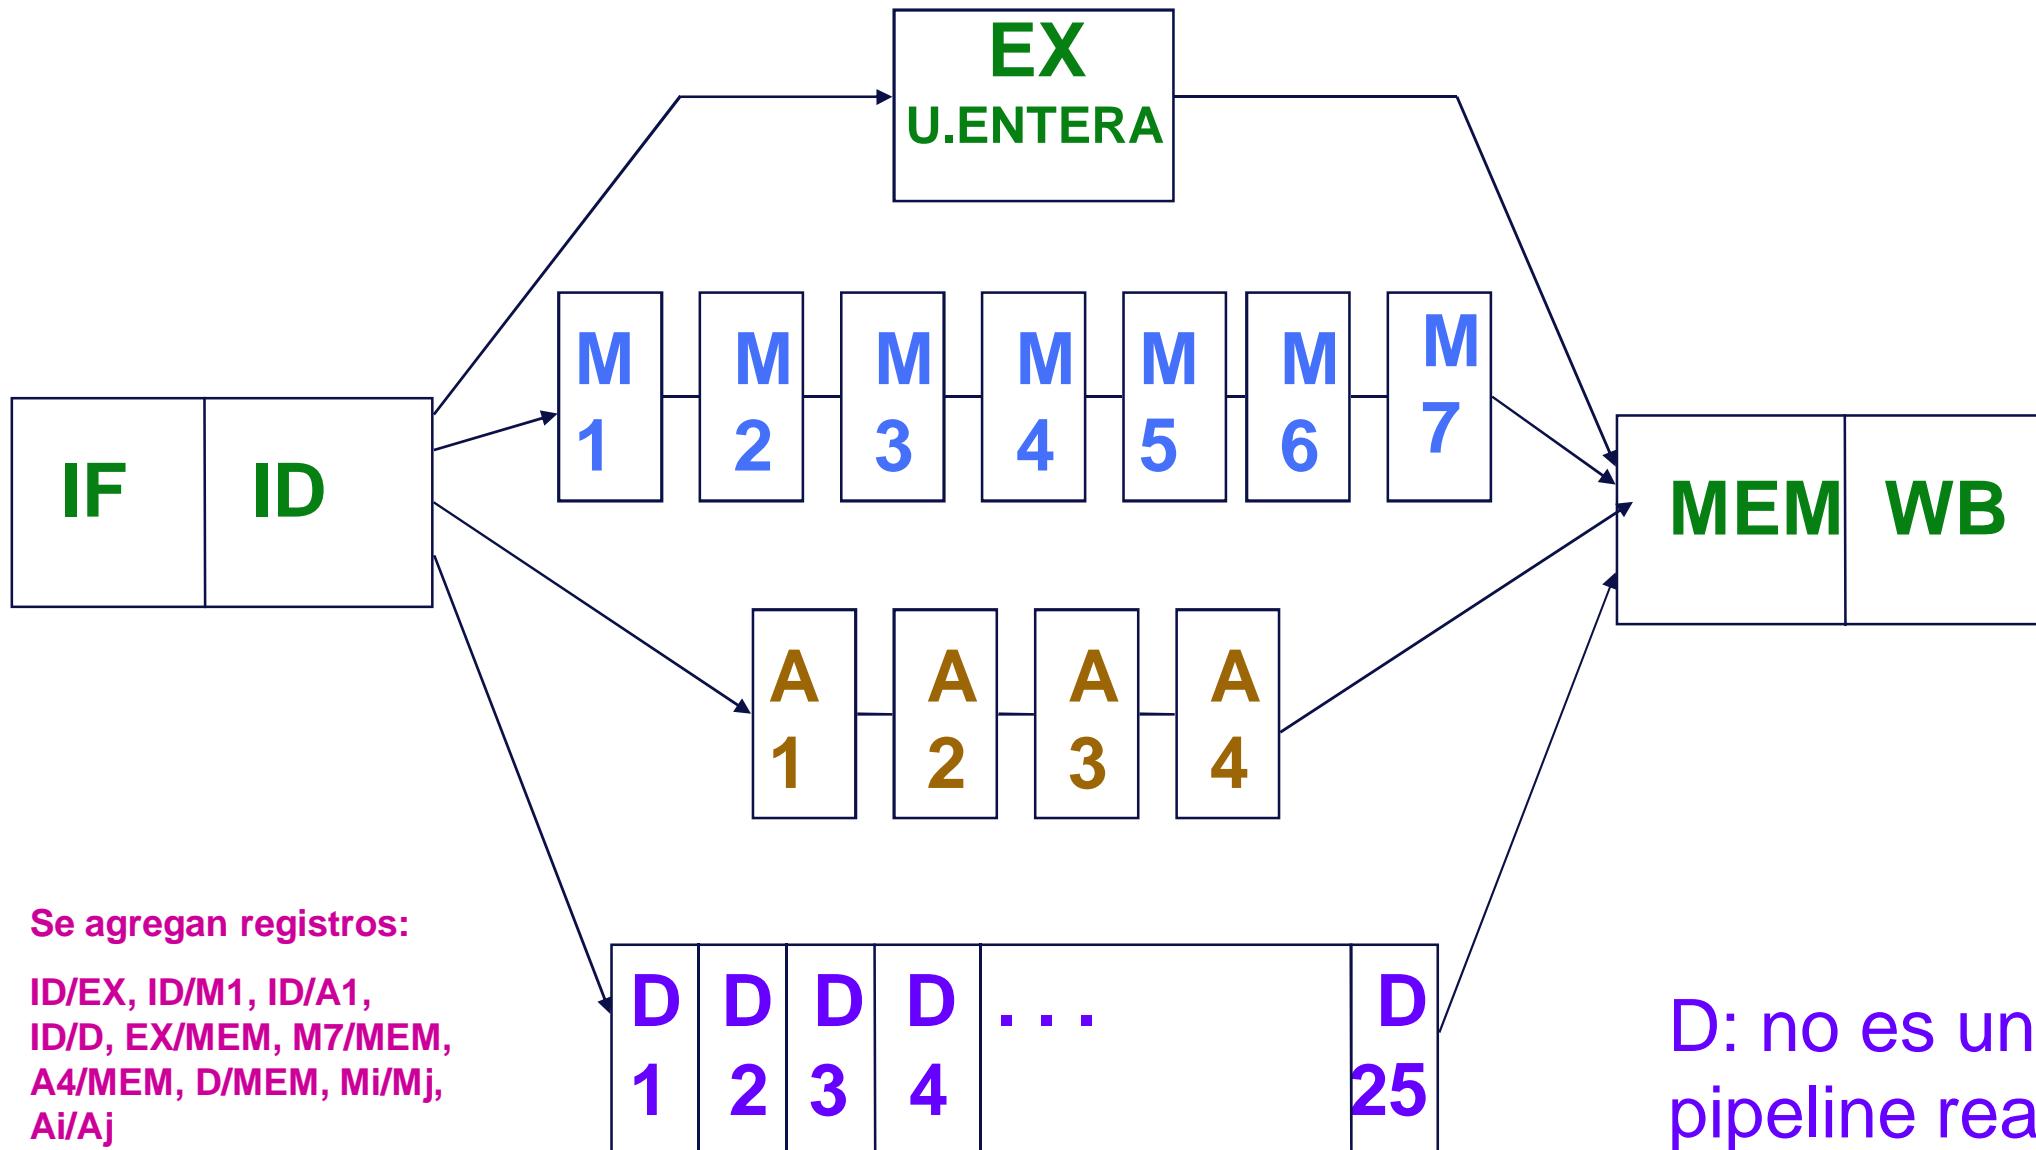


Figure C.16 Prediction accuracy of a 4096-entry 2-bit prediction buffer for the SPEC89 benchmarks. The misprediction rate for the integer benchmarks (gcc, espresso, eqntott, and li) is substantially higher (average of 11%) than that for the floating-point programs (average of 4%). Omitting the floating-point kernels (nasa7, matrix300, and tomcatv) still yields a higher accuracy for the FP benchmarks than for the integer benchmarks. These data, as well as the rest of the data in this section, are taken from a branch-prediction study done using the IBM Power architecture and optimized code for that system. See Pan et al. (1992). Although these data are for an older version of a subset of the SPEC benchmarks, the newer benchmarks are larger and would show slightly worse behavior, especially for the integer benchmarks.

PIPELINE RISC-V (Fig. C.30)



LATENCIA vs. INTERVALO DE REPETICIÓN

Latencia de unidad funcional: Número de **ciclos** entre instrucción que **usa resultados** y la instrucción que **los produce**.

Intervalo de iniciación/ intervalo de repetición: número de **ciclos** que deben darse **entre la emisión** de instrucciones del mismo tipo.

UNIDAD FUNCIONAL	LATENCIA	INTERVALO INIC.
ALU Entera	0	1
Memoria datos (loads ent. y fp, 1 menos para latencia de store)	1	1
Sumador/Rest. FP	3	1
Multiplicador Ent./FP	6	1
Divisor Ent/FP	24	25

Hacer aquí:

Ejemplos y práctica flotantes

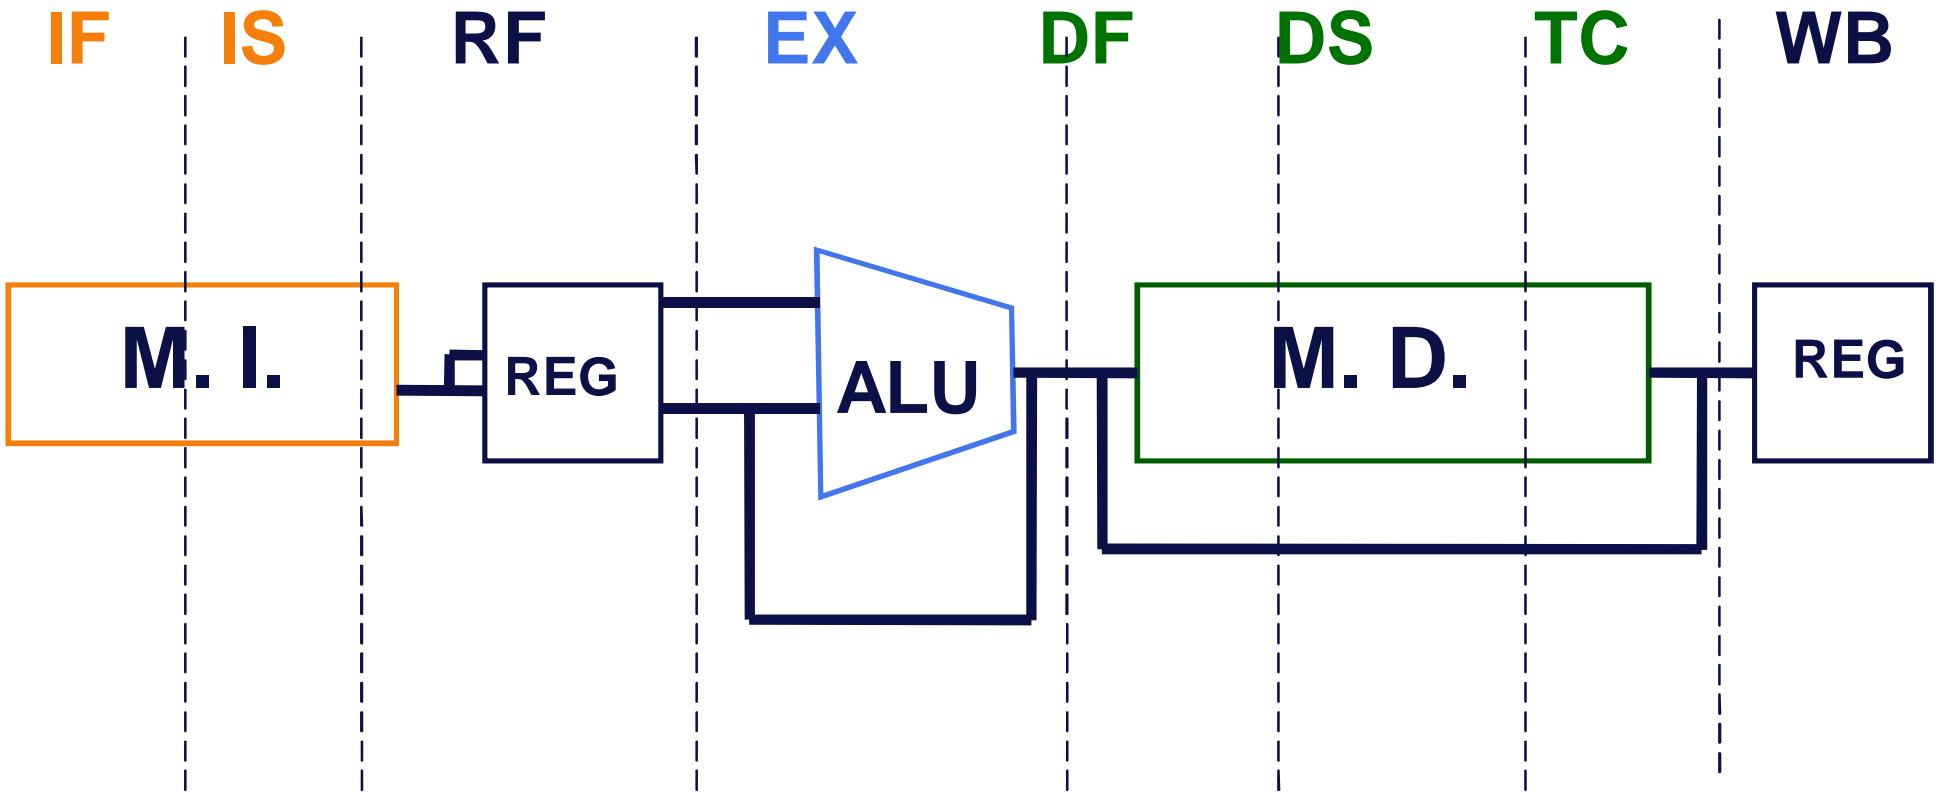
EL PIPELINE DE LA MIPS R4000 (A.6)

Es realmente una máquina de 64 bits

“Pipelining muy profundo (superpipelining). 8 etapas:

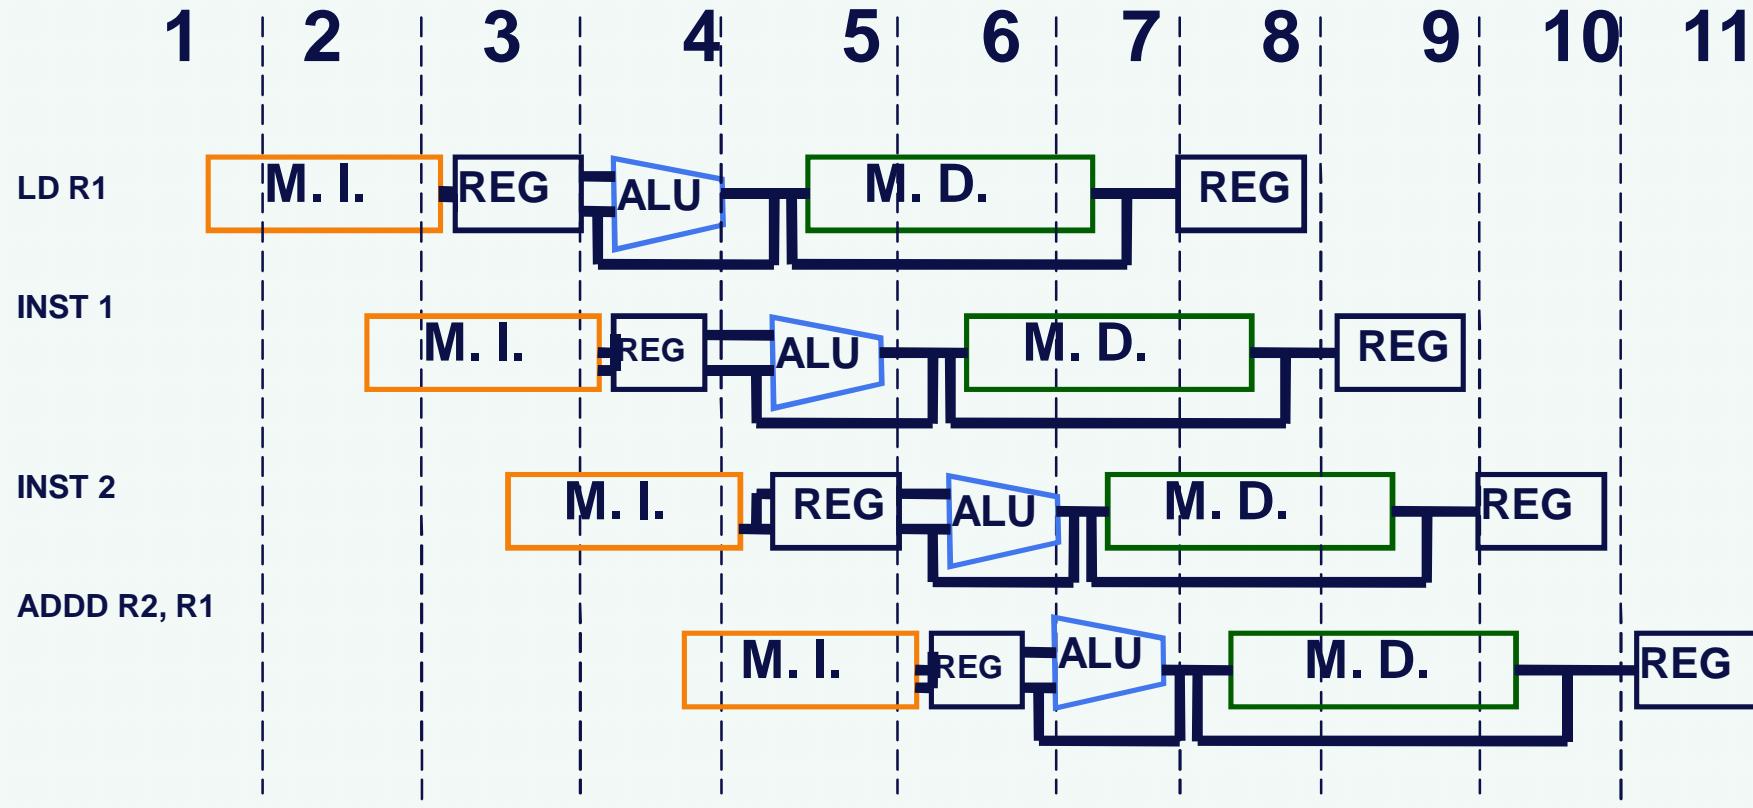
- “**IF** 1er parte fetch instrucc. selección PC, inicia inst. acceso a cache
- “**IS** 2da parte fetch instrucc. Se completa acceso a cache
- “**RF** Decod. Instruc, fetch registros, revisión conflictos, detección de si hubo hit en cache
- “**EX** EFA, ALU op, branch target (dirección nueva) y cód. condición
- “**DF** Fetch de datos, 1era mitad acceso a cache de datos
- “**DS** 2da mitad fetch de datos. Se completa acceso de datos en cache
- “**TC** Revisión de etiqueta (tag), se determina si hubo hit en acceso a cache de datos
- “**WB** Write Back+para operaciones reg-reg y loads

PIPELINE MIPS R4000 (CONT)



Load tiene 2 ciclos de retraso-(delay slots)

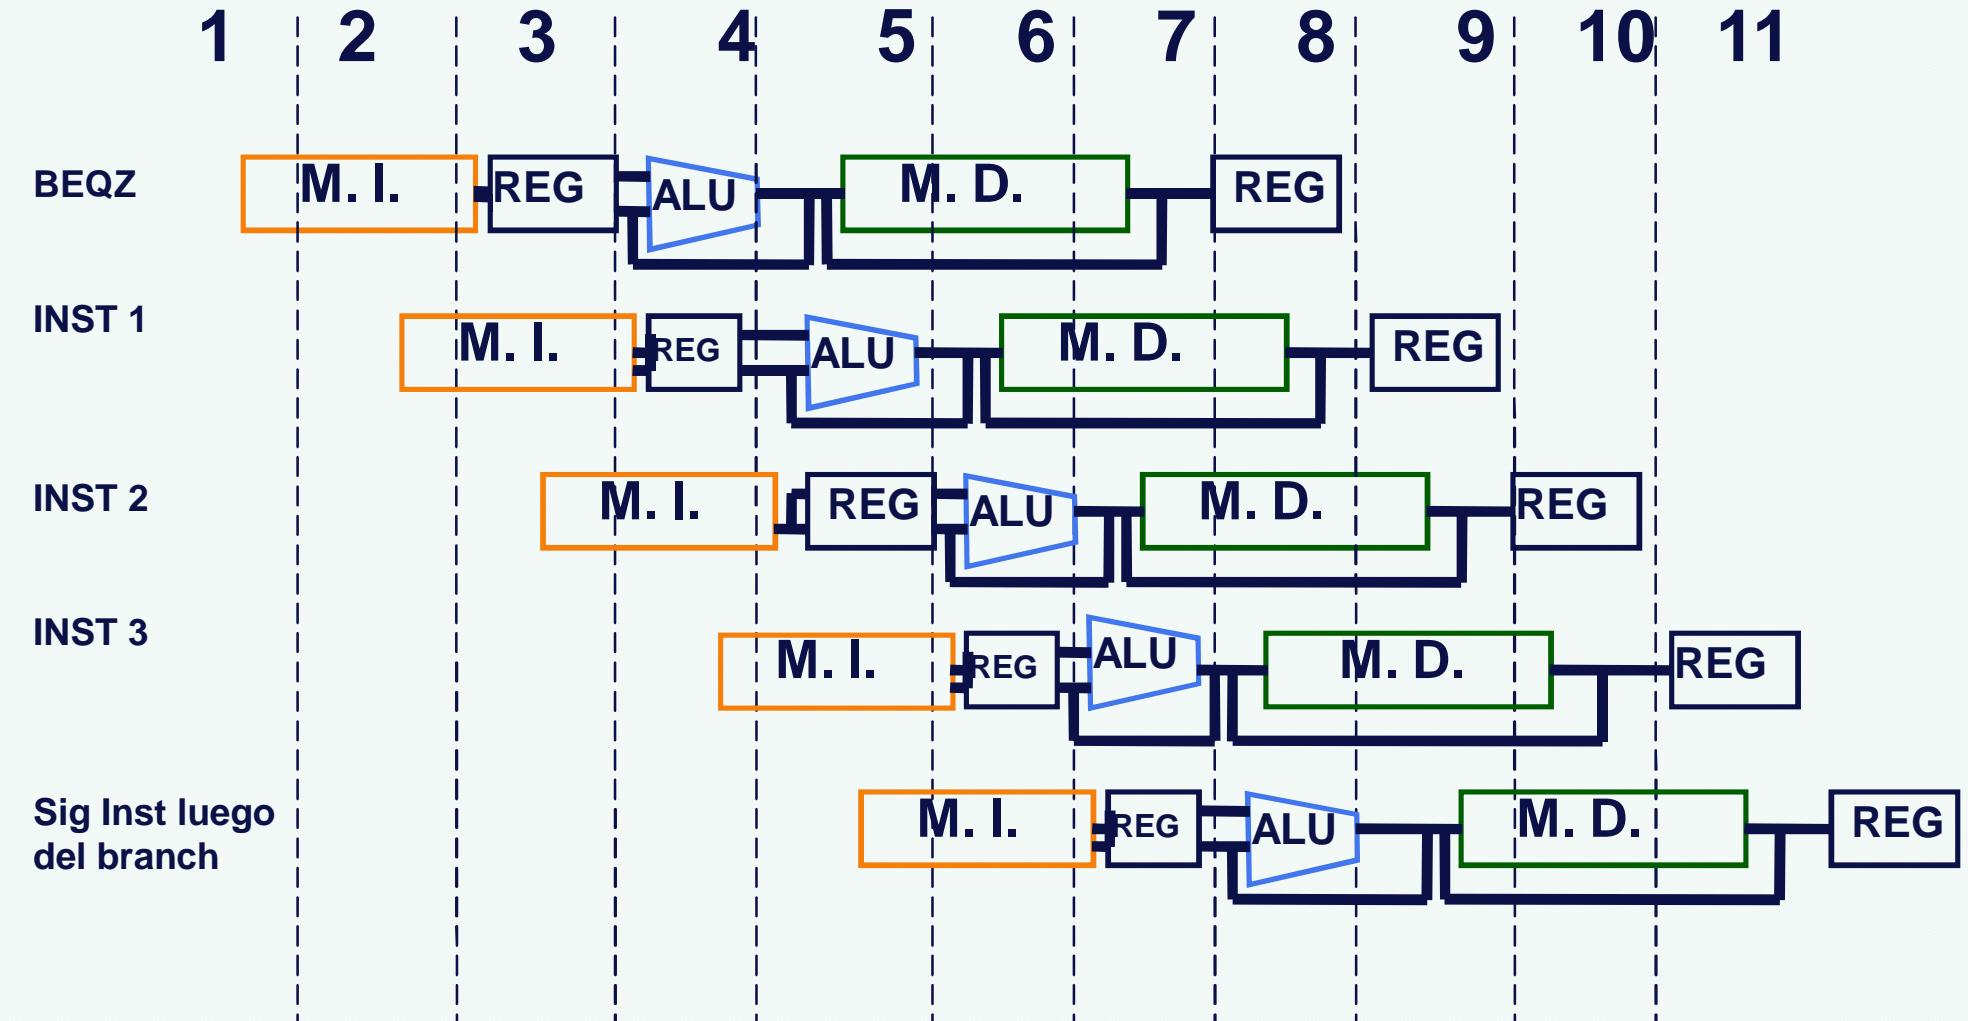
=> debe haber 2 instrucciones o burbujas entre un load y una instrucción dependiente de ésta.



Nótese que el valor está disponible al final de DS del load y se hace un **forwarding** al ALU del ADD=> se usa **ANTES** de la revisión de etiqueta del caché (tag check). Por lo que si hay una falta en cache hay que %devolverse+y atrasarse un ciclo. (ocurre menos del 10% de las veces y es fácil detener instrucciones dependientes)

BRANCH TIENE 3 CICLOS DE RETRASO (se resuelve en EX) =>
 CPI = 4 si no hay estrategia para disminuir retraso

(Con branch retrasado se pueden incluir 3 instrucciones luego del branch:)



MIPS SIEMPRE INCLUYE UNA INSTRUCCIÓN POR BRANCH RETRASADO
y procesador hace predicción no tomado para las siguientes instrucciones:

(# 1 cuando Branch fue tomado, # 2 cuando Branch no fue tomado)

# 1	1	2	3	4	5	6	7	8	9
BRANCH	IF	IS	RF	EX	DF	DS	TC	WB	
DELAY SLOT		IF	IS	RF	EX	DF	DS	TC	WB
BRANCH + 2			IF	ID					
BRANCH + 3				IF					
B.TARGET					IF	IS	RF	EX	DF

# 2	1	2	3	4	5	6	7	8	9 ...
BRANCH	IF	IS	RF	EX	DF	DS	TC	WB	
DELAY SLOT		IF	IS	RF	EX	DF	DS	TC	WB
BRANCH + 2			IF	IS	RF	EX	DF	DS	TC ...
BRANCH + 3				IF	IS	RF	EX	DF	DS ...
BRANCH + 4					IF IS	RF	EX	DF ...	

PIPELINE R4000 PARA PUNTO FLOTANTE

8 TIPOS DE ETAPAS

3 unidades funcionales: FP-DIV, FP-MUL, FP-ADD

Un sumador usado en paso final de MUL y DIV

ETAPA	UNIDAD FUNCIONAL	DESCRIPCION
A	FP-ADD	Suma mantisa
D	FP-DIV	Etapa de división
E	FP-MUL	Prueba de Excepción
M	FP-MUL	Primera etapa Multipl.
N	FP-MUL	Segunda etapa Multipl.
R	FP-ADD	Redondeo
S	FP-ADD	Í ShiftÍ de Operando
U	Las 3 unid. Func.	Í UnpackÍ número FP

OPERACIONES FP EN LA R4000

INST. FP	LATENCIA	INTERVALO INICIACION	SECUENCIA ETAPAS EN EL PIPE
Suma, resta	4	3	U, S + A, A + R, R + S
Multiplicación	8	4	U, E + M, M, M, M, N, N + A, R
División	36	35	U, A, R, D ²⁷ , D + A, D + R D + R, D + A, D + R, A, R
Raíz Cuadrada	112	111	U, E, (A + R) ¹⁰⁸ , A, R
Negación	2	1	U, S
Valor Absoluto	2	1	U, S
Comparación FP	3	2	U, A, R

Instruction type/opcode	Instruction meaning
<i>Data transfers</i>	
lb, lbu, sb	Move data between registers and memory, or between the integer and FP or special registers; only memory address mode is 12-bit displacement+contents of a GPR
lh, lhu, sh	Load byte, load byte unsigned, store byte (to/from integer registers)
lw, lwu, sw	Load half word, load half word unsigned, store half word (to/from integer registers)
ld, sd	Load word, load word unsigned, store word (to/from integer registers)
f1w, f1d, fsw, fsd	Load double word, store double word (to/from integer registers)
fmv._.x, fmv.x._	Load SP float, load DP float, store SP float, store DP float
fmv._.x, fmv.x._	Copy from/to integer register to/from floating-point register; “_”=S for single-precision, D for double-precision
csrrw, csrrwi, csrrs, csrrsi, csrrc, csrrci	Read counters and write status registers, which include counters: clock cycles, time, instructions retired
<i>Arithmetic/logical</i>	<i>Operations on integer or logical data in GPRs</i>
add, addi, addw, addiw	Add, add immediate (all immediates are 12 bits), add 32-bits only & sign-extend to 64 bits, add immediate 32-bits only
sub, subw	Subtract, subtract 32-bits only
mul, mulw, mulh, mulhsu, mulhu	Multiply, multiply 32-bits only, multiply upper half, multiply upper half signed-unsigned, multiply upper half unsigned
div, divu, rem, remu	Divide, divide unsigned, remainder, remainder unsigned
divw, divuw, remw, remuw	Divide and remainder: as previously, but divide only lower 32-bits, producing 32-bit sign-extended result
and, andi	And, and immediate
or, ori, xor, xori	Or, or immediate, exclusive or, exclusive or immediate
lui	Load upper immediate; loads bits 31-12 of register with immediate, then sign-extends
auipc	Adds immediate in bits 31–12 with zeros in lower bits to PC; used with JALR to transfer control to any 32-bit address
sll, slli, srl, srli, sra, srai	Shifts: shift left logical, right logical, right arithmetic; both variable and immediate forms
sllw, slliw, sr1w, srliw, sraw, sraiw	Shifts: as previously, but shift lower 32-bits, producing 32-bit sign-extended result
slt, slti, sltu, sltiu	Set less than, set less than immediate, signed and unsigned
<i>Control</i>	<i>Conditional branches and jumps; PC-relative or through register</i>
beq, bne, blt, bge, bltu, bgeu	Branch GPR equal/not equal; less than; greater than or equal, signed and unsigned
jal, jalr	Jump and link: save PC + 4, target is PC-relative (JAL) or a register (JALR); if specify ×0 as destination register, then acts as a simple jump
ecall	Make a request to the supporting execution environment, which is usually an OS
ebreak	Debuggers used to cause control to be transferred back to a debugging environment
fence, fence.i	Synchronize threads to guarantee ordering of memory accesses; synchronize instructions and data for stores to instruction memory

Instruction type/opcode	Instruction meaning
<i>Floating point</i>	<i>FP operations on DP and SP formats</i>
fadd.d, fadd.s	Add DP, SP numbers
fsub.d, fsub.s	Subtract DP, SP numbers
fmul.d, fmul.s	Multiply DP, SP floating point
fmadd.d, fmadd.s, fnmadd.d, fnmadd.s	Multiply-add DP, SP numbers; negative multiply-add DP, SP numbers
fmsub.d, fmsub.s, fnmsub.d, fnmsub.s	Multiply-sub DP, SP numbers; negative multiply-sub DP, SP numbers
fdiv.d, fdiv.s	Divide DP, SP floating point
fsqrt.d, fsqrt.s	Square root DP, SP floating point
fmax.d, fmax.s, fmin.d, fmin.s	Maximum and minimum DP, SP floating point
fcvt._._, fcvt._._u, fcvt._u._	Convert instructions: FCVT.x.y converts from type x to type y, where x and y are L (64-bit integer), W (32-bit integer), D (DP), or S (SP). Integers can be unsigned (U)
feq._, flt._, fle._	Floating-point compare between floating-point registers and record the Boolean result in integer register; “_”=S for single-precision, D for double-precision
fclass.d, fclass.s	Writes to integer register a 10-bit mask that indicates the class of the floating-point number ($-\infty$, $+\infty$, -0 , $+0$, NaN, ...)
fsgnj._, fsgnjn._, fsgnjx._	Sign-injection instructions that changes only the sign bit: copy sign bit from other source, the opposite of sign bit of other source, XOR of the 2 sign bits

Fig 1.6 Algunas operaciones para operandos de punto flotante de RISC-V

Program	Loads	Stores	Branches	Jumps	ALU operations
astar	28%	6%	18%	2%	46%
bzip	20%	7%	11%	1%	54%
gcc	17%	23%	20%	4%	36%
gobmk	21%	12%	14%	2%	50%
h264ref	33%	14%	5%	2%	45%
hmmer	28%	9%	17%	0%	46%
libquantum	16%	6%	29%	0%	48%
mcf	35%	11%	24%	1%	29%
omnetpp	23%	15%	17%	7%	31%
perlbench	25%	14%	15%	7%	39%
sjeng	19%	7%	15%	3%	56%
xalancbmk	30%	8%	27%	3%	31%

Figure A.29 RISC-V dynamic instruction mix for the SPECint2006 programs. Omnetpp includes 7% of the instructions that are floating point loads, stores, operations, or compares; no other program includes even 1% of other instruction types. A change in gcc in SPECint2006, creates an anomaly in behavior. Typical integer programs have load frequencies that are 1/5 to 3x the store frequency. In gcc, the store frequency is actually higher than the load frequency! This arises because a large fraction of the execution time is spent in a loop that clears memory by storing x0 (not where a compiler like gcc would usually spend most of its execution time!). A store instruction that stores a register pair, which some other RISC ISAs have included, would address this issue.